

Criteria for bias avoidance in machine learning

Examples of implicit bias abound in machine learning (ML) and statistics:

- Beauty filters make features more European (whiter skin, narrower noses). Sources: [the verge](#), [TechCrunch](#).
- An Google automated captioning system labeled Black people as "gorillas". Source: [Forbes](#)
- Word embeddings produce analogies such as "man is to software engineer as woman is to homemaker". Source: [Bolukbasi et al 2016](#)
- Crime prediction tools give higher false positive rates on Black people. [Source: ProPublica](#). Needless to say, this is a controversial topic. I haven't read it all and thought it through; that may be a future post. The ProPublica analysis is [here](#) and there is a rebuttal from the original authors [here](#).
- Polygenic scores (an application of machine learning to genetics and medicine) are expected to exacerbate health disparities. Source: [Martin et al arXiv preprint](#).

Motivated by my [wellness rewards post](#) ~~???TODO: FIX LINK~~, I began digging into the technical side of this problem: is it possible to do better, and if so, then how? In science, and especially statistical science, the devil is in the details. Even for the simplest prediction tasks, outcomes can vary drastically depending on your quantitative definition of fairness.

There are many mutually incompatible options, which are beginning to be laid out in textbooks and review papers. These typically avoid making value judgements about criteria, and that may serve their purposes, for example by allowing individual analysts to tailor solutions to their problems. But, that's often not necessary, because many of these problems follow the same basic format: machine learning classifiers are used to assign rewards or penalties to individuals. For this simple, widespread set of situations, there's really only one satisfactory criterion: the "counterfactual fairness" proposed [here](#) by Kusner et al.

Kusner, M. J., Loftus, J., Russell, C., & Silva, R. (2017). Counterfactual fairness. In Advances in Neural Information Processing Systems (pp. 4066-4076).

What problems does this apply to?

The scope of this post is narrow: I consider supervised machine learning problems with a single outcome of interest, where predictions directly affect a reward or penalty someone is offered. Examples might include predicting law school performance from pre-admission surveys (reward: admission), predicting crime rates by location (penalty: heavy policing), predicting job ad clicks based on user data (reward: access to high-paying jobs), or credit risk prediction (reward: a loan).

Some examples match the statistical format -- supervised ML with a single outcome -- but don't really fit what I have in mind because the resulting decision does not offer anyone a reward or penalty. Discrimination arises differently in such situations, and different fairness criteria are needed. An example is [clinical use of polygenic risk scores](#), where it might make sense to use an *equal accuracy* criterion that I argue against below.

(Un?)fairness through unawareness

- Definition: predictions are fair if they do not explicitly use protected factors A .
- Example: grocery-based wellness rewards prediction could be based on grocery receipts, without directly using race or gender.

Unfortunately, this does not work. The issue is that protected factors often explain part of the link between non-protected features X and the outcome Y . For instance, grocery purchases carry information about race, and health-oriented eating scores stratify people by race. A predictive method using grocery bills for health insurance pricing is likely to correlate with race whether you mean it to or not.

Setting various things equal within groups

- Definition: To achieve a criterion called *equal calibration*, predicted rates are required to match actual rates within each stratum. For all possible values of a ,

$$Pr(\hat{Y} = 1|A = a) = Pr(Y = 1|A = a)$$

$$E(\hat{Y}|A = a) = E(Y|A = a)$$

- Example: wellness rewards prediction meets this criterion if, under the program, health insurance prices for White people match health insurance costs for White people, and so on.

This seems like a minor variation on the strategy above: instead of maximizing an accuracy metric that averages over unobserved protected attributes, you could do this by maximizing accuracy separately within each category.

Unlike the unawareness strategy, predictions optimized for equal calibration will explicitly use information about protected attributes. Unfortunately, the information will be used to perpetuate historical inequities rather than to remedy them. For example, if Black clients have historically been more likely to default on loans, then risk prediction systems using the equal calibration criterion will result in lower access to capital for Black clients.

Auto insurance price discrimination by gender is commonplace in the US. People justify higher prices for men by pointing to higher rates of car crashes, car crash deaths, speeding, drunk driving, and not using seat belts. This is essentially an equal calibration argument. This is legal in all U.S. states except Montana ([source](#)) and, as of very recently, California ([source](#)). Here in Massachusetts, gender identity is a protected attribute for many

purposes, but pricing car insurance is not one of them ([source](#)). My interpretation is that equal calibration makes sense for some purposes, but it's got nothing to do with protecting against discrimination, and in fact it would be better called a justification for discrimination.

Setting various things equal across groups

Equal accuracy

- Definition: It is often possible to adjust predictions until some performance metric is the same across groups. I call this family of criteria the *equal accuracy family*. For example, one could require raw accuracy (for binary outcomes) or mean squared errors (for quantitative outcomes) to be the same for men and women.

$$Pr(\hat{Y} = Y|A = w) = Pr(\hat{Y} = Y|A = m)$$

$$E[(\hat{Y} - Y)^2|A = w] = E[(\hat{Y} - Y)^2|A = m]$$

- Example:

Prediction parity

- Definition: A similar option (sometimes called *statistical parity* or *demographic parity*, but I prefer the name *prediction parity*) is to require the distribution of predictions within each group to have the same mean.

$$Pr(\hat{Y} = 1|A = w) = Pr(\hat{Y} = 1|A = m)$$

I've only seen this applied to binary outcomes, and it could generalize to quantitative outcomes in various ways, which I won't discuss much. One obvious possibility is this.

$$E[\hat{Y}|A = w] = E[\hat{Y}|A = m]$$

- Example:

Equality of opportunity

- Definition: This is a variant of prediction parity or equal accuracy where equality is only required in part of the population.
- Example: consider recidivism prediction for men and women. Suppose Y is coded as 0 for people who don't recidivate (don't commit another crime after getting out of jail). Then equality of opportunity might be stated as equality of false positive rates.

$$Pr(\hat{Y} = 1 | Y = 0, A = w) = Pr(\hat{Y} = 1 | Y = 0, A = m)$$

The only difference is that the metric ignores what happens when $Y = 1$. If you use this metric, you're saying that for people who recidivate, it's okay for the prediction system to do a better job catching some than others, but among people who don't commit further crimes, men and women deserve equal treatment.

This criterion only makes sense for special situations. The use in law enforcement is problematic because one mechanism of racial discrimination by police is actually different rates of leniency towards *guilty* people, and because [in some prediction systems, the outcome is merely arrest and booking, not conviction by a jury of peers](#).

Ratings for criteria that set things equal across groups

In general, it's not possible to reconcile prediction parity with any type of equal accuracy criterion. This family of similar-but-incompatible metrics struggles to make sense of situations in which base rates differ between groups. If you plan to use a criterion from this family, you need a solid argument for which one it should be.

In many situations, existing outcomes Y were shaped by systems of advantage based on protected classes, so maintaining equal accuracy will not produce fair results. In fact, equal accuracy seems better suited to justifying discrimination rather than preventing it; I am hard pressed to come up with an intuitive example where equal accuracy works. Prediction parity might be a better criterion to use, but it's a blunt hammer and it can be viewed as penalizing groups with lower base rates.

I give a 'C' rating to prediction parity and a 'D' rating to equal accuracy.

Counterfactual fairness

- Definition: A recent paper by James Kusner and colleagues (citation below) advocates for the following definition of "Counterfactual Fairness". For all possible values of x , a , and a' ,

$$Pr(\hat{Y} | X = x, A = a) = Pr(\hat{Y}_{A \leftarrow a'} | X = x, A = a)$$

.

In plain English, if you alter someone's gender identity or racial/ethnic background (for example), it should not affect the predictions made about them.

- Example:

This really threw me for a loop: what exactly does the intervention entail? For biological/physical sex, the intervention seems simple to imagine: at conception, replace the paternal X chromosome with the paternal

Y chromosome or vice versa. For a concept like race, which is tied up with culture, history, and perceptions of others, it's not at all clear to me what this intervention would mean.

A key aspect of this definition is that you have to allow the intervention's effects to propagate. For the case of biological/physical sex, the intervention happens at conception, and a fair prediction should be invariant to all the resulting changes in either life experience or innate characteristics. If you really wouldn't have recidivated but for the testosterone, that counts as protected according to this criterion.

Regarding race, Kusner et al specify that the set of protected characteristics needs to be ancestrally closed with respect to the underlying causal graph. If race is protected, so is "mother's race", and so on. This has massive implications when applying the criterion to race as a protected attribute. If you want your prediction system to meet Kusner et al's standard of racial equity, you should imagine the intervention happening generations ago, with all the [cascading experiences of racial identity](#), and your predictions should still come out the same.

Here's the paper if you want to dig into it yourself.

Kusner, M. J., Loftus, J., Russell, C., & Silva, R. (2017). Counterfactual fairness. In Advances in Neural Information Processing Systems (pp. 4066-4076). [link](#)

This criterion allows for a lot of nuance: the analyst can structure a causal graph, specifying which variables and pathways are protected. It is correspondingly very difficult to justify models and conduct inference. I view such difficulties as a positive feature, because an honest criterion ought to reflect the obvious truth that this type of problem is difficult or impossible to solve.

Optimizing society-wide recovery from injustice

The criteria above are individualistic and symmetrical. This limits their ability to express what we know about the history behind our problems: race relations and gender roles in the US are certainly not symmetrical, and historically they are even less so. It also means these criteria encode only a mandate to be fair to directly affected people, rather than a mandate to make society as a whole more just. This presents a disconnect with broader ideas of how the justice system should work ([restorative justice](#)) or how to achieve fair access to [capital](#) and [educational opportunities](#). Whenever possible, we should step outside the intellectual box provided by a specific prediction problem and look at the landscape afresh.x

Appendix 1: Related work

There is a substantial amount of work on these problems. The [algorithmic justice league](#) takes a broad view and focuses on organizing researchers, journalists, companies, and policymakers. Google has posted a [discussion of best practices](#). On the technical side, there's a great beginner's introduction at <https://fairmlbook.org/>, with more to come as the textbook progresses. Tools for data scientists are aggregated

in the [IBM AI Fairness Toolkit](#), and there is also ongoing work at [Microsoft](#).

[These papers](#) and <https://fairmlbook.org/> present a nice problem framework and an overview of fairness criteria. Here are the papers.

Mitchell, S., Potash, E., & Barocas, S. (2018). Prediction-Based Decisions and Fairness: A Catalogue of Choices, Assumptions, and Definitions. arXiv preprint arXiv:1811.07867.

Berk, R., Heidari, H., Jabbari, S., Kearns, M., & Roth, A. (2017). Fairness in criminal justice risk assessments: the state of the art. arXiv preprint arXiv:1703.09207.

Appendix 2: Some notation

- Y denotes the outcome to be predicted. For example, this might be car accident risk if you work at an insurance company or it might be some aspect of criminal record if you're working on predictive policing.
- X denotes the set of features to be used in making predictions. For example, this might include location of residence or income. X may be statistically related to A : for example, in America, neighborhoods are highly segregated by race.
- A denotes a protected characteristic such as race or gender identity. For example, $A = 0$ could encode someone who is transgender and $A = 1$ someone who is cisgender.
- \hat{Y} denotes a prediction of Y . It is a function of X and A .
- U denotes unobserved variables that may be correlated and/or causally linked with X , A , and Y .
- Conditional probability distributions will be written $Pr(Z|W = w)$. The intuitive meaning of this is that you pick pairs Z, W at random and discard them unless $W = w$. For example, to sample from $Pr(\text{horsepower}|\text{carcolor} = \text{red})$, you pick cars until you get a red one, then measure its horsepower. For another example, you sample from $Pr(\text{interview}|\text{CandidateName} = \text{StereotypicallyBlack})$ by flipping through resumés until you get one with a stereotypically Black name, and then you figure out whether your company called them for an interview.
- Counterfactual probability distributions will be written with an extra subscript: $Pr(Z_{W \leftarrow u} | W = w)$. The intuitive meaning of this is that you pick pairs Z, W at random and discard them unless $W = w$, but then there's an additional step where you intervene and set W to u . For example, to sample from $Pr(\text{horsepower}_{\text{carcolor} \leftarrow \text{black}} | \text{carcolor} = \text{red})$, you pick cars until you get a red one, you paint it black, and then you measure its horsepower. For another example, to sample from

$Pr(\text{interview}_{\text{CandidateName} \leftarrow \text{StereotypicallyWhite}} | \text{CandidateName} = \text{StereotypicallyBlack})$

, you pick a resumé as before, change the name, and figure out whether your company called them for an interview. (Similar studies [have been done](#) to rigorously assess racial preferences in hiring.)